

NBA Shot Quality & Optimization (SqO)

Atiksh Singh
University of Connecticut
 atiksh.singh@uconn.edu

May 2025

Abstract

Modern basketball analytics have reached an “efficiency frontier” where traditional metrics no longer provide a competitive advantage. This paper introduces Shot Quality & Optimization (SqO), a spatial analytics framework that models personalized shooting efficiency for individual NBA players. Instead of relying on league-wide averages or simple zone shot charts, SqO builds a continuous Gaussian Kernel Density Estimation (KDE) surface from real shot chart data ingested via the official NBA Stats API. The resulting Points Per Shot (PPS) heatmaps and zone-level Expected Value (EV) audit tables give coaching staffs and front offices a data-driven tool to identify each player’s high-value zones, exploit opponent weaknesses, and make more precise decisions in high-stakes playoff scenarios. A working prototype is demonstrated across three archetypal players: Stephen Curry (elite shooter), Giannis Antetokounmpo (rim attacker), and DeMar DeRozan (midrange specialist). We are using 4,169 real shot attempts from the 2023-24 NBA season.

1. Introduction

The “Moneyball” era transformed basketball analytics, making concepts like corner three efficiency and points-per-possession mainstream. Yet as every team adopted the same playbook, the strategic advantage eroded. The advantage now is not what shots are good on average, it is which shots are good for a given player in a given context. As we know different NBA players have different assignments, a star player would have heavier expectations on them, given more opportunities to score, etc. than a role player. This kind of variation needs to be considered for strategy, something that league wide averages cannot give you.

Traditional shot quality models suffer from two key limitations. First, they rely on discrete zone classifications (e.g., K-Means clustering into 5-7 court regions) that smooth over meaningful spatial variation within zones. Second, they apply league-wide averages rather than learning individual player shooting signatures. A corner three is worth 1.11 expected points per shot

(PPS) for the average NBA player. For Giannis Antetokounmpo in 2023-24, it was worth 0.00 PPS, a number that changes how every defense in the league should guard him.

This paper introduces SqO (Shot Quality & Optimization), a prototype analytics tool that generates continuous, player-personalized PPS surfaces using Gaussian KDE smoothing on real NBA shot chart data. The system ingests data directly from the official `nba_api`, computes spatially-aware efficiency surfaces, and outputs both visual heatmaps and quantitative zone-level EV audits suitable for coaching and front-office decision making.

2. Related Work

2.0.1. K-Means Shot Zone Clustering

The most common industry baseline groups shots into discrete zones (restricted area, paint, midrange, corner 3, above-break 3) and computes FG% per zone. While

interpretable, this approach cannot capture variation within zones and treats all players identically within a zone classification.

2.0.2. Expected Possession Value (EPV)

Cervone et al. (2014) introduced EPV, a continuous model that estimates the expected point value of every moment in a possession using player-tracking data [2]. This is the most sophisticated production model, requiring SportVU or Second Spectrum tracking streams. SqO targets a practically accessible middle ground: richer than zone clustering, deployable without complicated tracking hardware.

2.0.3. CourtVision and Spatial Analytics

Goldsberry (2012) pioneered court-spatial visualization [1], building the foundation for court-as-grid modeling. SqO extends this by computing a statistically principled probability surface rather than binned frequency counts, and applies it to personalized player modeling.

2.0.4. Kernel Density Estimation in Sports

KDE has been applied to player positioning heatmaps in soccer and hockey [3]. Its application to shot quality (PPS surfaces) rather than shot frequency is a distinguishing contribution of this work.

3. Solution Design

3.1. Data Pipeline

Shot chart data is ingested via the `nba_api` Python library, which wraps the official NBA Stats API endpoints. For each player, the `ShotChartDetail` endpoint returns all field goal attempts for a specified season, including court coordinates (x, y) in tenths of feet from the basket, shot distance in feet, shot type (2PT or 3PT), and a made/missed flag.

Coordinates are normalized to feet: $x_{ft} = \frac{x}{10}$ and $y_{ft} = \frac{y}{10}$. A three-point indicator is assigned geometrically: a shot is classified as a three-pointer if its Euclidean distance from the basket is greater than or equal to 23.75 ft, or if $|x| \geq 22$ ft (corner three geometry), consistent with official NBA court specifications.

3.2. KDE Statistical Model

SqO uses Gaussian KDE smoothing to convert discrete shot events into a continuous court surface. The core insight is that each shot attempt provides information not just about the exact (x, y) location, but also about nearby regions and how much nearby information is used depends on the smoothing parameter σ .

For a 50×50 grid across the half-court, the smoothed make and attempt counts at grid cell (i, j) are:

$$\tilde{m}_{ij} = \sum_k m_k \cdot G(d_{ijk}, \sigma)$$

$$\tilde{a}_{ij} = \sum_k G(d_{ijk}, \sigma)$$

where $G(d, \sigma)$ is a Gaussian kernel, $m_k \in \{0, 1\}$ is the make indicator for shot k , and d_{ijk} is the distance from shot k to grid cell (i, j) .

A Bayesian prior of $w = 5$ pseudo-shots anchored to the league average FG% ($\mu_0 = 0.46$) is applied to prevent overfitting in sparse court zones:

$$\widehat{\text{FG}}_{0ij} = \frac{\tilde{m}_{ij} + w \cdot \mu_0}{\tilde{a}_{ij} + w}$$

The final output is a Points Per Shot surface that accounts for shot value:

$$\text{PPS}_{ij} = \widehat{\text{FG}}_{0ij} \times v_{ij}$$

where $v_{ij} = 3$ for three-point territory and $v_{ij} = 2$ otherwise.

3.3. Visualization and Reporting

The PPS surface is rendered as a heatmap using `matplotlib`, with color normalized independently within each scoring zone (inside and outside the arc). This normalization reveals efficiency zones that would otherwise be hidden by the large absolute value difference between two-point and three-point territory.

A zone-level EV audit divides each player's shots into four strategic zones: Rim (distance less than 6 ft), Corner Three ($|x| \geq 22$ ft), Midrange (2PT non-rim), and Above-Break Three (3PT non-corner), calculating the empirical EV per zone.

4. Solution Implementation

4.1. Prototype Stack

The prototype is implemented in Python using the following libraries: `nba_api` for data ingestion, `numpy` and `scipy.ndimage.gaussian_filter` for KDE computation, `pandas` for data manipulation, and `matplotlib` for visualization. No GPU compute or neural network training is required as the model runs to completion in under 60 seconds on standard CPU hardware.

4.2. Dataset

Three players were selected to represent 3 distinct shooting archetypes, allowing for direct comparison of how the model captures individual signatures:

Player	Archetype	Shots
Stephen Curry	Elite Shooter	1,445
Giannis Antet.	Rim Attacker	1,369
DeMar DeRozan	Midrange	1,355

All data is from the 2023-24 NBA regular season.

4.3. Results: Zone-Level EV Audit

The table below shows the empirical Expected Value (points per shot) for each player across four strategic court zones, alongside league average baselines.

Player	Rim	Cor. 3	Mid	AbvBrk 3
Curry	1.244	1.571	0.830	1.183
Giannis	1.472	0.000	0.690	0.850
DeRozan	1.227	1.167	0.916	0.894
Lg. Avg.	1.20	1.11	0.78	1.05

Three findings stand out with direct strategic analysis.

Finding 1 — Curry’s corner three (1.571 PPS): This is 41% above league average, the highest single-zone EV in the dataset. Defenses must treat Curry as a priority close-out target at all times in the corner, even without the ball.

Finding 2 — Giannis’s corner three (0.000 PPS): Zero made corner threes in 1,369 attempts. Defenses should hard-shade Giannis toward corner three posi-

tions on every possession — it is mathematically the lowest-EV shot he can take. This contradicts the conventional wisdom that corner threes are universally high-value.

Finding 3 — DeRozan’s midrange (0.916 PPS):

This is 17% above league average (0.78), in an era where the midrange shot is getting rarer as players are encouraged to leave the arc for their shots more often. DeRozan’s shooting signature demonstrates that player-personalized models are essential, league-average thinking would incorrectly flag his bread-and-butter shot as inefficient.

5. Lessons Learned and Beyond

5.1. Model Methodology Decision

The initial prototype plan was to use a deep neural network (TensorFlow MLP) to regress shot success probability from $(x, y, \text{distance})$ features. This approach produced two critical failure modes: first, training instability due to random weight initialization caused different outputs on identical data between runs for example different PPS, EV, and heatmaps, and second, high variance in sparse court zones where fewer than 10 shots were attempted.

Switching to KDE with a Bayesian prior resolved both issues. The model is fully deterministic, statistically principled for small-sample spatial data, and produces outputs that are easier to interpret and explain to coaching stakeholders. This aligns with established practices in the sports industry where KDE is the standard approach in spatial point process modeling [4].

5.2. Primary Challenges

Data sparsity: Some court zones have very few shot attempts, particularly in the corners for players who rarely shoot there. The Bayesian prior addresses this but does not eliminate it entirely for extreme cases.

Absence of defensive context: The current model uses only shot outcomes and locations. It cannot distinguish a wide-open three from a heavily contested one. Incorporating defender proximity data from Second Spectrum tracking would be the single highest-value improvement.

Static signatures: The model treats all shots from a season equally, ignoring fatigue, game-state context, and opponent-specific tendencies. A player’s efficiency in the fourth quarter of a playoff game may differ significantly from their regular-season average.

5.3. Future Roadmap

Phase 2 — Defender Proximity: Integrate Second Spectrum tracking data to condition efficiency estimates on defensive pressure, enabling separate surfaces for open versus contested shot scenarios.

Phase 3 — Game-State Conditioning: Use sequence modeling (LSTM or Transformer architecture) to estimate how shot quality changes as a function of game clock, score differential, and fatigue proxy metrics.

Phase 4 — Opponent Scouting Module: Generate defensive vulnerability maps by inverting the model — identifying which zones each defense surrenders most efficiently, enabling targeted offensive game-planning.

Phase 5 — Real-Time API Integration: Deploy as a REST API that ingests live play-by-play data and updates PPS mid-game, supporting in-game timeout decisions.

6. Conclusion

SqO demonstrates that player-personalized spatial analytics can reveal strategic insights not highlighted in traditional zone-based models. By applying Gaussian KDE smoothing to real NBA shot chart data, the prototype generates PPS surfaces that quantify each player’s shooting signature across the full court.

The zone-level EV audit surfaces three immediately actionable findings: Curry’s corner three is 41% more valuable than league average, Giannis’s corner three is statistically worthless, and DeRozan’s midrange defies the conventional efficiency narrative. These are the kinds of insights that inform playoff defensive schemes and shot selection decisions worth millions of dollars in competitive advantage.

7. References

1. Goldsberry, K. (2012). **CourtVision: New Visual Analytics for Consumer Basketball Data**. MIT Sloan Sports Analytics Conference.
2. Cervone, D., D’Amour, A., Bornn, L., & Goldsberry, K. (2014). **POINTWISE: Predicting Points and Valuing Decisions in Real Time**. MIT Sloan Sports Analytics Conference.
3. Franks, A., Miller, A., Bornn, L., & Goldsberry, K. (2015). **Counterpoints: Advanced Defensive Metrics for NBA Basketball**. MIT Sloan Sports Analytics Conference.
4. Diggle, P. J. (2013). **Statistical Analysis of Spatial and Spatio-Temporal Point Patterns** (3rd ed.). CRC Press.
5. Sexton, S. et al. (2024). **nba_api: An API Client for NBA Stats**. https://github.com/swar/nba_api